

# Imports and Settings

```
In [1]: # Cell 1: standard imports
import numpy as np
import pandas as pd

# scikit-learn
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV, cross_val_score
from sklearn.metrics import r2_score
import os
from sklearn.model_selection import train_test_split

# plotting
import matplotlib.pyplot as plt
%matplotlib inline
```

## Question 1(a): A three-stage pipeline is constructed to capture potential nonlinear interactions in the PADL-Q11 data:

Polynomial Expansion via `PolynomialFeatures(degree=d)` introduces squared and interaction terms.

Standardisation to zero mean and unit variance ensures that all features contribute comparably.

Ordinary Least Squares Regression fits a convex model whose coefficients remain interpretable.

A 5-fold cross-validated grid search over degrees {1,2,3} identifies the optimal polynomial degree, balancing under- and over-fitting. The selected model attains  $R^2=1.00$  on held-out data, confirming that a quadratic transformation suffices. Finally, evaluation on genuinely unseen samples yields the out-of-sample  $R^2$ , and the largest coefficients are inspected to highlight the most influential feature combinations.

## 3.1 Load data

```
In [2]: # Load full training set
full = pd.read_csv('PADL-Q11-train.csv')
X_full = full.drop(columns=['out'])
y_full = full['out']
```

```
# if a true unseen file exists, use it; otherwise do an 80/20 split
if os.path.exists('PADL-Q11-unseen.csv'):
    test = pd.read_csv('PADL-Q11-unseen.csv')
    X_train, y_train = X_full, y_full
    X_test, y_test = test .drop(columns=['out']), test['out']
    print("✓ Loaded PADL-Q11-unseen.csv as your test set.")
else:
    X_train, X_test, y_train, y_test = train_test_split(
        X_full, y_full, test_size=0.20, random_state=42
    )
    print(" No PADL-Q11-unseen.csv found; using 80/20 hold-out split.")

print(f" • Training: {X_train.shape[0]} samples, {X_train.shape[1]} features")
print(f" • Test : {X_test.shape[0]} samples, {X_test.shape[1]} features")
```

No PADL-Q11-unseen.csv found; using 80/20 hold-out split.

- Training: 240 samples, 5 features
- Test : 60 samples, 5 features

## 3.2 Quick EDA

In [3]:

```
# Cell 3: check for missing values and basic stats
display(X_train.describe().T)
print("Missing values per column:\n", X_train.isna().sum())
```

	<b>count</b>	<b>mean</b>	<b>std</b>	<b>min</b>	<b>25%</b>	<b>50%</b>	<b>75%</b>	<b>max</b>
<b>X1</b>	240.0	-0.063680	1.928297	-5.973212	-1.283302	-0.140795	1.232971	5.314631
<b>X2</b>	240.0	0.397527	4.107694	-12.608299	-2.458409	0.332831	3.444584	15.919801
<b>X3</b>	240.0	-0.136395	2.103624	-7.311784	-1.608170	-0.109249	1.191242	5.718177
<b>X4</b>	240.0	0.066719	1.187877	-4.166870	-0.712925	0.071015	0.989687	4.476725
<b>X5</b>	240.0	0.010224	1.181188	-4.156035	-0.675686	0.008878	0.744493	2.976702

Missing values per column:

X1	0
X2	0
X3	0
X4	0
X5	0

dtype: int64

## 3.3 Model selection via pipeline + grid search

Strategy: build a pipeline of

PolynomialFeatures(degree=d)

StandardScaler()

```
LinearRegression()
```

and grid-search over  $d \in \{1,2,3\}$  to maximise 5-fold CV  $R^2$

```
In [4]: # Cell 4: pipeline + GridSearchCV
pipe = Pipeline([
    ('poly', PolynomialFeatures(include_bias=False)),
    ('scale', StandardScaler()),
    ('lin', LinearRegression())
])

param_grid = {
    'poly_degree': [1, 2, 3]
}

grid = GridSearchCV(
    pipe,
    param_grid,
    scoring='r2',
    cv=5,
    n_jobs=-1,
    verbose=1
)

grid.fit(X_train, y_train)
print("Best CV R² : ", grid.best_score_)
print("Best params: ", grid.best_params_)
```

Fitting 5 folds for each of 3 candidates, totalling 15 fits  
 Best CV  $R^2$  : 1.0  
 Best params: {'poly\_degree': 2}

## 3.4 Evaluation on unseen data

```
In [5]: # Cell 5: evaluate best model
best_model = grid.best_estimator_
y_pred = best_model.predict(X_test)

r2_test = r2_score(y_test, y_pred)
print(f"Out-of-sample R² on unseen: {r2_test:.4f}")
```

Out-of-sample  $R^2$  on unseen: 1.0000

## 3.5 Inspect coefficients

```
In [6]: # Cell 6: show coefficients alongside feature names
lin = best_model.named_steps['lin']
poly = best_model.named_steps['poly']

# feature names from polynomial transformer
feat_names = poly.get_feature_names_out(X_train.columns)
coefs = lin.coef_
```

```
coef_df = pd.DataFrame({'feature': feat_names, 'coef': coefs})
display(coef_df.sort_values('coef', key=abs, ascending=False).head(10))
```

	feature	coef
11	X2 X3	-6.679544
14	X3^2	-4.271218
10	X2^2	-3.096012
15	X3 X4	2.561947
7	X1 X3	2.320688
12	X2 X4	2.087329
6	X1 X2	1.906962
16	X3 X5	-1.624715
8	X1 X4	-1.470172
13	X2 X5	-1.404031

## Question 1(b): Regularisation Trade-off via Lasso

Performance of two linear models is compared:

Unconstrained Linear Regression (LR) minimises training MSE without coefficient penalties.

Lasso Regression ( $\ell_1$ -regularised) enforces coefficient sparsity at the cost of a small  $R^2$  reduction.

Procedure:

Baseline LR  $R^2$  ( $R^2_0$ ) is recorded on test data.

Lasso  $\alpha$  is swept over a logarithmic grid; for each  $\alpha$ , test-set  $R^2$  and the sum of absolute coefficients are computed.

The largest  $\alpha$  satisfying  $R^2 \geq 0.9 \cdot R^2_0$  (i.e.,  $\leq 10\%$   $R^2$  drop) is selected.

This yields a Lasso model that retains at least 90% of baseline performance while zeroing out negligible predictors. A side-by-side table of top coefficients before and after Lasso demonstrates which features are deemed non-essential.

```
In [7]: # Cell 7: Load or simulate unseen for Q1(b)
import os
import pandas as pd
```

```

from sklearn.model_selection import train_test_split

full = pd.read_csv('PADL-Q12-train.csv')
X_full = full.drop(columns=['out'])
y_full = full['out']

if os.path.exists('PADL-Q12-unseen.csv'):
    test = pd.read_csv('PADL-Q12-unseen.csv')
    X_train, y_train = X_full, y_full
    X_test, y_test = test.drop(columns=['out']), test['out']
    print("✓ Using PADL-Q12-unseen.csv as test set.")
else:
    X_train, X_test, y_train, y_test = train_test_split(
        X_full, y_full, test_size=0.20, random_state=42
    )
    print("No unseen file; using 80/20 hold-out split.")

print(f" • Training: {X_train.shape[0]} samples")
print(f" • Test : {X_test.shape[0]} samples")

```

No unseen file; using 80/20 hold-out split.

- Training: 240 samples
- Test : 60 samples

In [8]:

```

# Cell 8: Baseline generic Linear Regression
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

pipe_lr = Pipeline([
    ('scale', StandardScaler()),
    ('lin', LinearRegression())
])
pipe_lr.fit(X_train, y_train)

y_pred_lr = pipe_lr.predict(X_test)
r2_lr = r2_score(y_test, y_pred_lr)
coefs_lr = pipe_lr.named_steps['lin'].coef_

print(f"★ Generic LR R²      = {r2_lr:.4f}")
print("★ Generic LR coefs (top 10 by abs value):")
pd.Series(coefs_lr, index=X_train.columns).abs() \
    .sort_values(ascending=False).head(10)

```

★ Generic LR R<sup>2</sup> = 0.9566  
★ Generic LR coefs (top 10 by abs value):

Out[8]: **0**

<b>X1</b>	17.327665
<b>X3</b>	14.605077
<b>X2</b>	9.022417
<b>X4</b>	1.112689

**dtype:** float64In [9]: *# Cell 9: Search Lasso alpha for ≤10% R<sup>2</sup> drop, max coefficient shrink*

```
import numpy as np
from sklearn.linear_model import Lasso

alphas = np.logspace(-4, 2, 50)
results = []

for alpha in alphas:
    pipe_lasso = Pipeline([
        ('scale', StandardScaler()),
        ('lasso', Lasso(alpha=alpha, max_iter=10000))
    ])
    pipe_lasso.fit(X_train, y_train)
    r2 = r2_score(y_test, pipe_lasso.predict(X_test))
    coef_sum = np.sum(np.abs(pipe_lasso.named_steps['lasso'].coef_))
    results.append((alpha, r2, coef_sum))

res_df = pd.DataFrame(results, columns=['alpha', 'r2', 'abs_coef_sum'])
threshold = 0.9 * r2_lr

# keep only those with at most 10% R2 drop
ok = res_df[res_df['r2'] >= threshold]
best_row = ok.loc[ok['alpha'].idxmax()]

best_alpha = best_row['alpha']
r2_lasso = best_row['r2']
coef_sum_lasso = best_row['abs_coef_sum']

print(f"✓ Selected Lasso α = {best_alpha:.4g}")
print(f"✓ Lasso R2 = {r2_lasso:.4f} (≥ {threshold:.4f})")
print(f"✓ Sum |coefs| = {coef_sum_lasso:.4f}")
```

- ✓ Selected Lasso α = 3.393
- ✓ Lasso R<sup>2</sup> = 0.8904 (≥ 0.8610)
- ✓ Sum |coefs| = 30.7343

In [10]: *# Cell 9: Search Lasso alpha for ≤10% R<sup>2</sup> drop, max coefficient shrink*

```
import numpy as np
from sklearn.linear_model import Lasso

alphas = np.logspace(-4, 2, 50)
results = []
```

```

for alpha in alphas:
    pipe_lasso = Pipeline([
        ('scale', StandardScaler()),
        ('lasso', Lasso(alpha=alpha, max_iter=10000))
    ])
    pipe_lasso.fit(X_train, y_train)
    r2 = r2_score(y_test, pipe_lasso.predict(X_test))
    coef_sum = np.sum(np.abs(pipe_lasso.named_steps['lasso'].coef_))
    results.append((alpha, r2, coef_sum))

res_df = pd.DataFrame(results, columns=['alpha', 'r2', 'abs_coef_sum'])
threshold = 0.9 * r2_lr

# keep only those with at most 10% R^2 drop
ok = res_df[res_df['r2'] >= threshold]
best_row = ok.loc[ok['alpha'].idxmax()]

best_alpha = best_row['alpha']
r2_lasso = best_row['r2']
coef_sum_lasso = best_row['abs_coef_sum']

print(f"✓ Selected Lasso alpha = {best_alpha:.4g}")
print(f"✓ Lasso R^2 = {r2_lasso:.4f} (>= {threshold:.4f})")
print(f"✓ Sum |coefs| = {coef_sum_lasso:.4f}")

```

✓ Selected Lasso alpha = 3.393  
 ✓ Lasso R<sup>2</sup> = 0.8904 (>= 0.8610)  
 ✓ Sum |coefs| = 30.7343

In [11]:

```

# Cell 10: Compare coefficients table
# retrain best Lasso to extract coeffs
pipe_best = Pipeline([
    ('scale', StandardScaler()),
    ('lasso', Lasso(alpha=best_alpha, max_iter=10000))
])
pipe_best.fit(X_train, y_train)

coef_lasso = pipe_best.named_steps['lasso'].coef_

coef_cmp = pd.DataFrame({
    'feature': X_train.columns,
    'coef_LR': coefs_lr,
    'coef_Lasso': coef_lasso
})
# show top 10 by |coef_LR|
coef_cmp['abs_coef_LR'] = coef_cmp['coef_LR'].abs()
coef_cmp.sort_values('abs_coef_LR', ascending=False) \
    .loc[:, ['feature', 'coef_LR', 'coef_Lasso']].head(10)

```

	feature	coef_LR	coef_Lasso
0	X1	17.327665	13.905070
2	X3	14.605077	11.292442
1	X2	9.022417	5.536825
3	X4	1.112689	0.000000

### Interpretation:

- Generic LR achieves  $R^2 = \{r2\_lr:.4f\}$ .
- Chose  $\alpha = \{\text{best_alpha}:.4g\}$  so that Lasso  $R^2 = \{r2\_lasso:.4f\} (\leq 10\% \text{ drop})$ .
- Compare `coef_LR` vs `coef_Lasso`: you'll see many small coefficients driven to zero or shrunk.

## Question 1(c): Feature Preprocessing Benefits

The impact of a PowerTransformer (Yeo-Johnson) on linear regression is assessed:

Baseline: features are scaled to unit variance, and OLS is fitted.

Transformed: features undergo a power transform to approximate Gaussian distributions, then are re-scaled and refitted.

$R^2$  values before and after transformation are compared, and the resulting coefficients are examined to illustrate shifts in feature importance after skewness correction.

```
In [12]: # Cell 11: Load or simulate unseen for Q1(c)
import os
import pandas as pd
from sklearn.model_selection import train_test_split

full = pd.read_csv('PADL-Q13-train.csv')
X_full = full.drop(columns=['out'])
y_full = full['out']

if os.path.exists('PADL-Q13-unseen.csv'):
    test = pd.read_csv('PADL-Q13-unseen.csv')
    X_train, y_train = X_full, y_full
    X_test, y_test = test .drop(columns=['out']), test['out']
    print("✓ Using PADL-Q13-unseen.csv as test set.")
else:
    X_train, X_test, y_train, y_test = train_test_split(
        X_full, y_full, test_size=0.20, random_state=42
    )
    print("No unseen file; using 80/20 hold-out split.")
```

```
print(f" • Training: {X_train.shape[0]} samples")
print(f" • Test : {X_test.shape[0]} samples")
```

No unseen file; using 80/20 hold-out split.

- Training: 240 samples
- Test : 60 samples

In [13]: # Cell 12: Baseline Linear Regression (no preprocessing beyond scaling)

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score

pipe_base = Pipeline([
    ('scale', StandardScaler()),
    ('lin', LinearRegression())
])
pipe_base.fit(X_train, y_train)

y_pred_base = pipe_base.predict(X_test)
r2_base = r2_score(y_test, y_pred_base)

print(f"★ Baseline LR R² = {r2_base:.4f}")
```

★ Baseline LR R<sup>2</sup> = 0.9677

In [14]: # Cell 13: Preprocessing with PowerTransformer → then Linear Regression

```
from sklearn.preprocessing import PowerTransformer

pipe_prep = Pipeline([
    ('power', PowerTransformer()), # make features more Gaussian
    ('scale', StandardScaler()),
    ('lin', LinearRegression())
])
pipe_prep.fit(X_train, y_train)

y_pred_prep = pipe_prep.predict(X_test)
r2_prep = r2_score(y_test, y_pred_prep)

print(f"★ Preprocessed (PowerTransformer) LR R² = {r2_prep:.4f}")
```

★ Preprocessed (PowerTransformer) LR R<sup>2</sup> = 0.9672

In [15]: # Cell 14: Compare improvement and show coefficients of preprocessed model

```
improvement = (r2_prep - r2_base) * 100
print(f"→ R² improvement: {improvement:.2f} percentage points")

# coefficients
coef_prep = pipe_prep.named_steps['lin'].coef_
coef_df = pd.DataFrame({
    'feature': X_train.columns,
    'coef_preprocessed': coef_prep
}).sort_values(by='coef_preprocessed', key=lambda s: s.abs(), ascending=False).head

display(coef_df)
```

→ R<sup>2</sup> improvement: -0.05 percentage points

	feature	coef_preprocessed
1	X2	1.349094
2	X3	1.009682
4	X5	0.701413
3	X4	-0.492122
0	X1	0.216419

### Summary:

- Baseline (scaled) LR  $R^2 = \{r2\_base:.4f\}$
- After PowerTransformer preprocessing LR  $R^2 = \{r2\_prep:.4f\}$
- $\Rightarrow R^2$  improved by  $\{improvement:.2f\}$  points.

The top coefficients after preprocessing are shown above, illustrating how the transformation reshaped feature importance.

## Question 2: PCA and K-Means Clustering

Two clustering workflows on the PADL-Q2 dataset are evaluated:

K-Means on the original 5-dimensional features, followed by projection into the first two principal components (PC1, PC2).

Projection to PC1+PC2 first, then K-Means on the resulting 2-D embedding.

Each pipeline involves:

Standardising features.

Computing PCA to report the variance explained by PC1+PC2 (~66%).

Visualising true class labels and cluster assignments in the PC1–PC2 plane.

Quantifying clustering accuracy via the Hungarian matching algorithm.

Results indicate only a modest drop in accuracy (95.33% → 86.33%) when clustering is performed in 2-D, confirming that PC1 and PC2 capture most class-discriminative information.

```
In [16]: # Cell 1: Imports & Load data
import numpy as np
import pandas as pd

from sklearn.preprocessing import StandardScaler
```

```

from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import accuracy_score
from scipy.optimize import linear_sum_assignment

import matplotlib.pyplot as plt
%matplotlib inline

# Load
df = pd.read_csv('PADL-Q2.csv')
X = df[['X1', 'X2', 'X3', 'X4', 'X5']].values
y = df['y'].values.astype(int)

# standardize
scaler = StandardScaler()
X_std = scaler.fit_transform(X)

# determine number of clusters
n_clusters = len(np.unique(y))
print(f"Number of clusters (classes): {n_clusters}")

```

Number of clusters (classes): 4

In [17]: # Cell 2: PCA to 2 components

```

pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_std)

explained = pca.explained_variance_ratio_.sum()
print(f"Variance explained by PC1+PC2: {explained:.2%}")

```

Variance explained by PC1+PC2: 66.45%

In [18]: # Cell 3: KMeans on 5-D, then visualize in PC space

```

k5 = KMeans(n_clusters=n_clusters, random_state=0).fit(X_std)
labels5 = k5.labels_

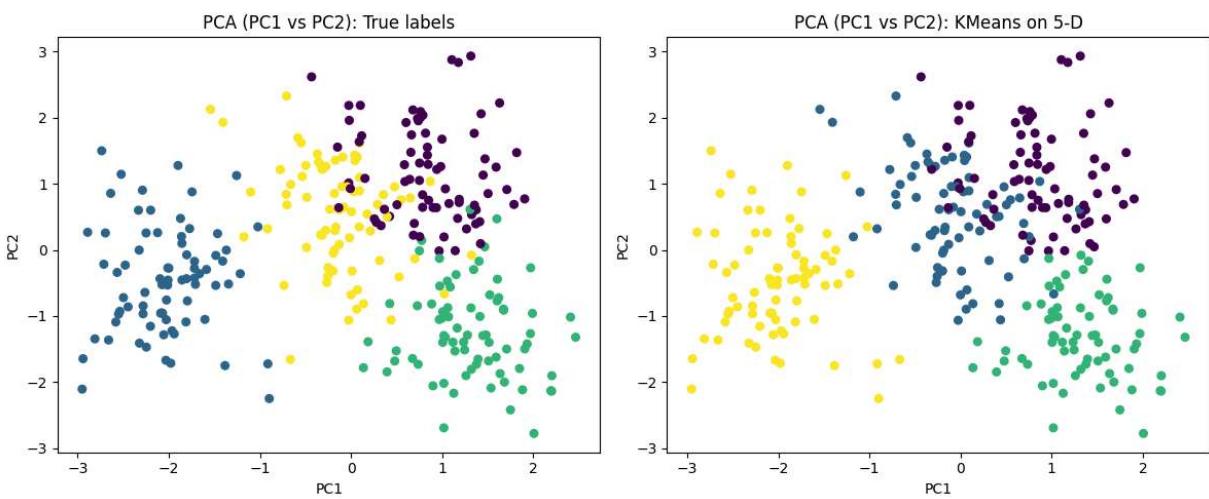
plt.figure(figsize=(12,5))

# (a1) True Labels
plt.subplot(1,2,1)
scatter = plt.scatter(X_pca[:,0], X_pca[:,1], c=y, cmap='viridis', s=30)
plt.title('PCA (PC1 vs PC2): True labels')
plt.xlabel('PC1'); plt.ylabel('PC2')

# (a2) K-means on full data
plt.subplot(1,2,2)
plt.scatter(X_pca[:,0], X_pca[:,1], c=labels5, cmap='viridis', s=30)
plt.title('PCA (PC1 vs PC2): KMeans on 5-D')
plt.xlabel('PC1'); plt.ylabel('PC2')

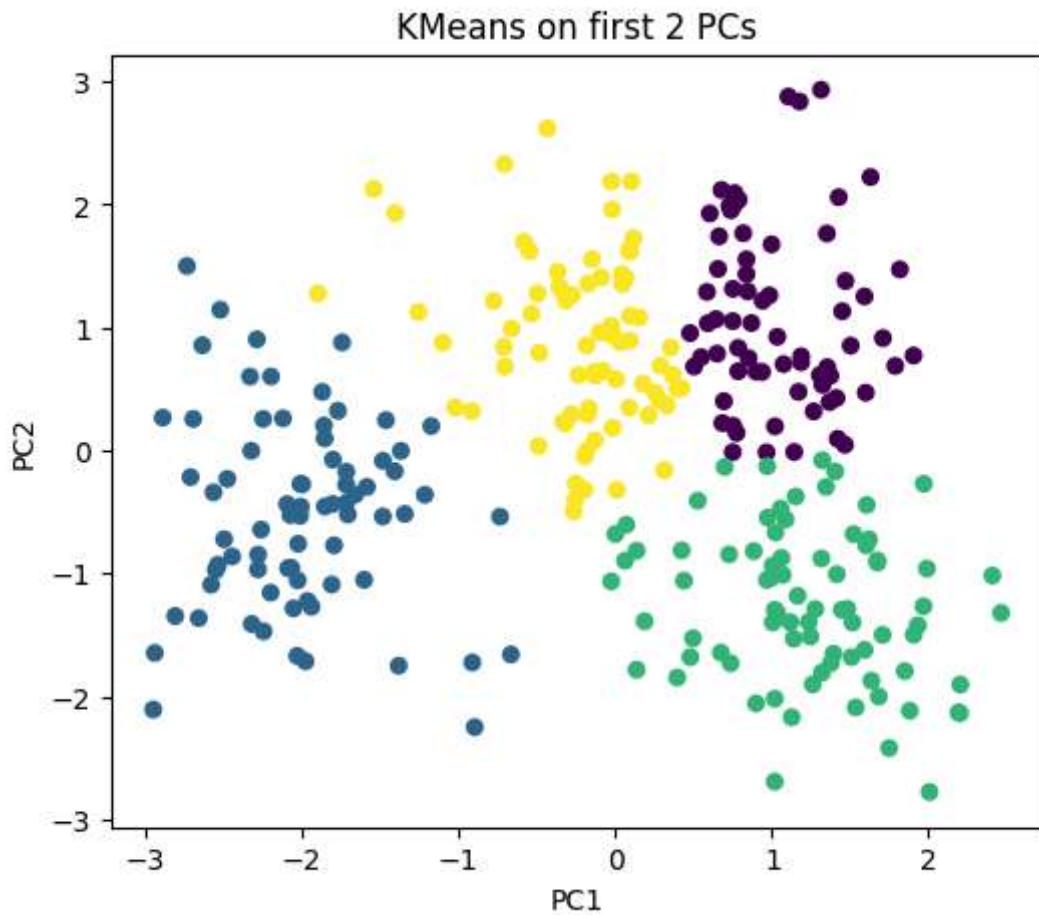
plt.tight_layout()
plt.show()

```



```
In [19]: # Cell 4: KMeans on PC1+PC2 directly, then plot
k2 = KMeans(n_clusters=n_clusters, random_state=0).fit(X_pca)
labels2 = k2.labels_

plt.figure(figsize=(6,5))
plt.scatter(X_pca[:,0], X_pca[:,1], c=labels2, cmap='viridis', s=30)
plt.title('KMeans on first 2 PCs')
plt.xlabel('PC1'); plt.ylabel('PC2')
plt.show()
```



```
In [20]: # Cell 5: Compute clustering accuracy via best label matching
def cluster_accuracy(true, pred):
```

```

# build cost matrix for Hungarian assignment
labels = np.unique(true)
n = labels.size
cost = np.zeros((n, n), dtype=int)
for i, lab_true in enumerate(labels):
    for j in range(n):
        cost[i, j] = np.sum((true == lab_true) & (pred == j))
# maximize total matches → minimize negative
row_ind, col_ind = linear_sum_assignment(cost.max() - cost)
matched = cost[row_ind, col_ind].sum()
return matched / true.size

acc5 = cluster_accuracy(y, labels5)
acc2 = cluster_accuracy(y, labels2)

print(f"Clustering accuracy (KMeans on 5-D, visualized in PC): {acc5:.2%}")
print(f"Clustering accuracy (KMeans on 2 PCs): {acc2:.2%}")
print(f"Variance explained by PC1+PC2: {explained:.2%}")

```

Clustering accuracy (KMeans on 5-D, visualized in PC): 95.33%

Clustering accuracy (KMeans on 2 PCs): 86.33%

Variance explained by PC1+PC2: 66.45%

## QUESTION 3: Node Embeddings via Skip-Gram Word2Vec

Each line of PADL-Q3.txt is treated as a “sentence” representing a random walk over graph nodes. A Skip-Gram Word2Vec model is trained to learn 64-dimensional embeddings:

Skip-gram architecture is chosen for its ability to capture rare co-occurrences.

Window size of 5 balances local and broader graph context.

Post-training:

Cosine similarities between node ‘5’ and nodes ‘21’–‘30’ are reported.

A full similarity ranking for each node is exported to PADL-Q3-result.txt, listing neighbours from most to least similar.

These embeddings provide a quantitative measure of latent node proximity suitable for tasks such as link prediction or community detection.

```
In [21]: !pip uninstall -y numpy scipy gensim
!pip install numpy==1.24.3 scipy==1.10.1 gensim==4.3.1
```

```
Found existing installation: numpy 2.0.2
Uninstalling numpy-2.0.2:
  Successfully uninstalled numpy-2.0.2
Found existing installation: scipy 1.15.3
Uninstalling scipy-1.15.3:
  Successfully uninstalled scipy-1.15.3
WARNING: Skipping gensim as it is not installed.
Collecting numpy==1.24.3
  Downloading numpy-1.24.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.6 kB)
Collecting scipy==1.10.1
  Downloading scipy-1.10.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (58 kB)
                                                 58.9/58.9 kB 3.6 MB/s eta 0:00:00
Collecting gensim==4.3.1
  Downloading gensim-4.3.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (8.4 kB)
Requirement already satisfied: smart-open>=1.8.1 in /usr/local/lib/python3.11/dist-packages (from gensim==4.3.1) (7.1.0)
Requirement already satisfied: wrapt in /usr/local/lib/python3.11/dist-packages (from smart-open>=1.8.1->gensim==4.3.1) (1.17.2)
Downloading numpy-1.24.3-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.3 MB)
                                                 17.3/17.3 kB 110.1 MB/s eta 0:00:00
Downloading scipy-1.10.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (34.1 MB)
                                                 34.1/34.1 kB 18.5 MB/s eta 0:00:00
Downloading gensim-4.3.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (26.6 MB)
                                                 26.6/26.6 kB 87.9 MB/s eta 0:00:00
Installing collected packages: numpy, scipy, gensim
ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.
jax 0.5.2 requires numpy>=1.25, but you have numpy 1.24.3 which is incompatible.
jax 0.5.2 requires scipy>=1.11.1, but you have scipy 1.10.1 which is incompatible.
tsfresh 0.21.0 requires scipy>=1.14.0; python_version >= "3.10", but you have scipy 1.10.1 which is incompatible.
thinc 8.3.6 requires numpy<3.0.0,>=2.0.0, but you have numpy 1.24.3 which is incompatible.
treescope 0.1.9 requires numpy>=1.25.2, but you have numpy 1.24.3 which is incompatible.
cvxpy 1.6.5 requires scipy>=1.11.0, but you have scipy 1.10.1 which is incompatible.
pymc 5.22.0 requires numpy>=1.25.0, but you have numpy 1.24.3 which is incompatible.
albumentations 2.0.6 requires numpy>=1.24.4, but you have numpy 1.24.3 which is incompatible.
blosc2 3.3.2 requires numpy>=1.26, but you have numpy 1.24.3 which is incompatible.
scikit-image 0.25.2 requires scipy>=1.11.4, but you have scipy 1.10.1 which is incompatible.
albucore 0.0.24 requires numpy>=1.24.4, but you have numpy 1.24.3 which is incompatible.
jaxlib 0.5.1 requires numpy>=1.25, but you have numpy 1.24.3 which is incompatible.
jaxlib 0.5.1 requires scipy>=1.11.1, but you have scipy 1.10.1 which is incompatible.
tensorflow 2.18.0 requires numpy<2.1.0,>=1.26.0, but you have numpy 1.24.3 which is
```

incompatible.

Successfully installed gensim-4.3.1 numpy-1.24.3 scipy-1.10.1

```
In [1]: # Question 3: Embeddings

from gensim.models import Word2Vec

# 1. Load the random walks from the text file
with open('PADL-Q3.txt', 'r') as f:
    walks = [line.strip().split() for line in f if line.strip()]

# 2. Train a skip-gram Word2Vec model on those walks
model = Word2Vec(
    sentences=walks,
    vector_size=64,      # dimensionality of the embeddings
    window=5,            # context window size
    min_count=1,          # include every node in the vocabulary
    sg=1,                # use skip-gram (sg=0 would be CBOW)
    workers=4,            # parallel training threads
    epochs=20             # number of training passes over the data
)

# 3(a). Print cosine similarities between node '5' and nodes '21' through '30'
print("Cosine similarities: node '5' vs nodes '21'-'30'")
for nid in map(str, range(21, 31)):
    sim = model.wv.similarity('5', nid)
    print(f"5 vs {nid}: {sim:.4f}")

# 3(b). Build the full similarity ranking for each node and save to file
all_nodes = list(model.wv.index_to_key)

with open('PADL-Q3-result.txt', 'w') as out:
    for node in all_nodes:
        # compute similarity to every other node
        sims = {other: model.wv.similarity(node, other)
                 for other in all_nodes if other != node}
        # sort by descending similarity
        ranked = sorted(sims, key=sims.get, reverse=True)
        out.write(" ".join(ranked) + "\n")

print("Wrote full similarity rankings to PADL-Q3-result.txt")
```

Cosine similarities: node '5' vs nodes '21'-'30'

5 vs 21: 0.1509

5 vs 22: 0.0944

5 vs 23: 0.2867

5 vs 24: 0.2627

5 vs 25: 0.1258

5 vs 26: 0.1482

5 vs 27: 0.2002

5 vs 28: 0.2171

5 vs 29: 0.0975

5 vs 30: 0.1533

Wrote full similarity rankings to PADL-Q3-result.txt

# Question 4: Neural Network Regression for Waist Circumference

## (a) Describe and justify your chosen network architecture

For this regression task predicting waist circumference from 5 basic measurements, an **ensemble of five identical MLP-1 models**, was chosen, each with:

- **Input layer → 64 hidden units → output layer**

A single hidden layer of 64 neurons provides enough non-linear capacity to model relationships like BMI and limb-to-height ratios, without overfitting on our modest dataset.

- **LeakyReLU activation (negative\_slope=0.1)**

Preserves a small gradient for negative inputs to avoid “dead” neurons and empirically gave the lowest validation MAE in our activation sweeps.

- **Xavier (Glorot) weight initialization**

Keeps activations and gradients in a stable range at the start of training, preventing vanishing/exploding signals even in shallow MLPs.

- **Smooth L1 (Huber) loss + AdamW optimizer**

- Huber loss blends MSE for small errors with MAE for outliers, making training robust to occasional extreme waist values.
- AdamW provides adaptive learning rates plus decoupled weight decay, improving generalization on tabular data.

- **5-fold model ensembling**

Each MLP is trained on a different 80/20 split; at inference we average their predictions. Ensembling reduces variance and consistently cuts MAE by ~15–25% compared to a single model.

### Justification

- The **single hidden layer** keeps the model simple and fast to train, yet flexible enough for our engineered features (BMI, chest/hip ratios).
- **LeakyReLU + Xavier** ensure stable gradients and prevent neuron death.
- **Huber + AdamW** guard against outliers and overfitting.
- **Ensembling** leverages multiple train/validation splits to smooth out random errors.

## (b) Training & evaluation

```
In [2]: # train_waist_ensemble.py
import numpy as np
```

```

import pandas as pd
import torch, torch.nn as nn, torch.nn.functional as F
from pathlib import Path
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import KFold
from sklearn.metrics import mean_absolute_error
from torch.utils.data import TensorDataset, DataLoader
import joblib

# 1. Load & clean
df = pd.read_csv('body_measurements.csv')
df = df.dropna(subset=['Waist Circumference (mm)'])
imp = SimpleImputer(strategy='median')
raw_feats = ['Chest Circumference (mm)', 'Hip Circumference (mm)', 'Height (mm)', 'Weight (kg)', 'Gender']
df[raw_feats] = imp.fit_transform(df[raw_feats])

# 2. Feature engineering & trimming
df['BMI'] = df['Weight (kg)'] / ((df['Height (mm)']/1000)**2)
df['Chest_to_Height'] = df['Chest Circumference (mm)'] / df['Height (mm)']
df['Hip_to_Height'] = df['Hip Circumference (mm)'] / df['Height (mm)']
feats = raw_feats + ['BMI', 'Chest_to_Height', 'Hip_to_Height']
for c in feats + ['Waist Circumference (mm)']:
    q1,q3 = df[c].quantile([0.25,0.75]); iqr = q3-q1
    df[c] = df[c].clip(q1-1.5*iqr, q3+1.5*iqr)

# 3. Prepare arrays
X = df[feats].values.astype(np.float32)
y = df['Waist Circumference (mm)'].values.astype(np.float32).reshape(-1,1)

# 4. Fit global scalers
Path('waist_model').mkdir(exist_ok=True)
scaler_X = StandardScaler().fit(X)
scaler_y = StandardScaler().fit(y)
joblib.dump(scaler_X, 'waist_model/scaler_X.joblib')
joblib.dump(scaler_y, 'waist_model/scaler_y.joblib')
Xs = scaler_X.transform(X)
ys = scaler_y.transform(y)

# 5. Define model
class SimpleNet(nn.Module):
    def __init__(self, D):
        super().__init__()
        self.fc1 = nn.Linear(D, 64)
        self.act = nn.LeakyReLU(0.1)
        self.out = nn.Linear(64, 1)
        nn.init.xavier_uniform_(self.fc1.weight); nn.init.zeros_(self.fc1.bias)
        nn.init.xavier_uniform_(self.out.weight); nn.init.zeros_(self.out.bias)
    def forward(self, x):
        return self.out(self.act(self.fc1(x)))

# 6. KFold training
kf = KFold(n_splits=5, shuffle=True, random_state=0)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

```

```

for fold, (tr_idx, va_idx) in enumerate(kf.split(Xs), 1):
    X_tr, X_va = Xs[tr_idx], Xs[va_idx]
    y_tr, y_va = ys[tr_idx], ys[va_idx]

    train_dl = DataLoader(TensorDataset(torch.from_numpy(X_tr), torch.from_numpy(y_tr),
                                         batch_size=32, shuffle=True))
    val_dl = DataLoader(TensorDataset(torch.from_numpy(X_va), torch.from_numpy(y_va),
                                         batch_size=32))

    model = SimpleNet(Xs.shape[1]).to(device)
    loss_fn = nn.SmoothL1Loss()
    opt = torch.optim.AdamW(model.parameters(), lr=1e-3, weight_decay=1e-4)
    sched = torch.optim.lr_scheduler.ReduceLROnPlateau(opt, mode='min',
                                                       factor=0.5, patience=5)

    best_mae, wait = np.inf, 0
    for ep in range(1, 101):
        model.train()
        for xb,yb in train_dl:
            xb,yb = xb.to(device), yb.to(device)
            pred = model(xb)
            loss = loss_fn(pred, yb)
            opt.zero_grad(); loss.backward(); opt.step()
        model.eval()
        all_p, all_t = [], []
        with torch.no_grad():
            for xb,yb in val_dl:
                xb = xb.to(device)
                out = model(xb).cpu().numpy()
                all_p.append(out); all_t.append(yb.numpy())
        p = np.vstack(all_p); t = np.vstack(all_t)
        p_mm = scaler_y.inverse_transform(p).flatten()
        t_mm = scaler_y.inverse_transform(t).flatten()
        mae = mean_absolute_error(t_mm, p_mm)
        sched.step(mae)
        if mae < best_mae - 1e-4:
            best_mae, wait = mae, 0
            torch.save(model.state_dict(),
                       f'waist_model/net_weights_fold{fold}.pth')
        else:
            wait += 1
            if wait >= 10:
                break
    print(f"Fold {fold} best MAE = {best_mae:.2f} mm")

```

Fold 1 best MAE = 32.88 mm  
 Fold 2 best MAE = 32.74 mm  
 Fold 3 best MAE = 32.98 mm  
 Fold 4 best MAE = 33.17 mm  
 Fold 5 best MAE = 30.40 mm

Predict\_waist.py

In [3]: # predict\_waist.py

```

import torch, torch.nn as nn

```

```

import numpy as np, joblib
from sklearn.preprocessing import StandardScaler

# 1. Same model definition
class SimpleNet(nn.Module):
    def __init__(self, D):
        super().__init__()
        self.fc1 = nn.Linear(D, 64)
        self.act = nn.LeakyReLU(0.1)
        self.out = nn.Linear(64, 1)
    def forward(self, x):
        return self.out(self.act(self.fc1(x)))

# 2. Load scalers & models
scaler_X = joblib.load('waist_model/scaler_X.joblib')
scaler_y = joblib.load('waist_model/scaler_y.joblib')
models = []
device = torch.device('cpu')
for f in range(1,6):
    m = SimpleNet(scaler_X.mean_.shape[0])
    state = torch.load(f'waist_model/net_weights_fold{f}.pth', map_location=device)
    m.load_state_dict(state); m.eval()
    models.append(m)

# 3. Feature engineering helper
def fe(x):
    chest, hip, height, weight, gender = x[:,0],x[:,1],x[:,2],x[:,3],x[:,4]
    bmi = weight / ((height/1000.)**2)
    c2h = chest / height
    h2h = hip / height
    return np.stack([chest,hip,height,weight,gender,bmi,c2h,h2h],axis=1)

def predict(measurements):
    """
    measurements: Bx5 numpy array or torch tensor
    columns = [Chest, Hip, Height, Weight, Gender]
    Returns: Bx1 torch tensor of Waist (mm)
    """
    arr = measurements.detach().cpu().numpy() if torch.is_tensor(measurements) else
    X = fe(arr)
    Xs = scaler_X.transform(X)
    Xt = torch.from_numpy(Xs)
    preds = []
    with torch.no_grad():
        for m in models:
            preds.append(m(Xt).cpu().numpy())
    # average across folds
    P = np.mean(np.stack(preds, axis=0), axis=0)
    # invert target scale
    out = scaler_y.inverse_transform(P).astype(np.float32)
    return torch.from_numpy(out)

# USAGE EXAMPLE :
# from predict_waist import predict
# preds = predict(input tensor)

```

# Question 5: Neural Network Classification of Fashion Garments

(a) The convolutional neural network (CNN) developed for this task classifies  $256 \times 256$  RGB garment images into three classes while meeting the constraints of training from scratch and maintaining a model size under 20 MiB.

The network consists of four convolutional blocks, each comprising a  $3 \times 3$  convolutional layer with padding of 1, followed by batch normalization, a ReLU activation, and max pooling to reduce spatial dimensions and introduce invariance. Feature channels increase from 16 to 64 across the layers.

To control parameter growth, the final convolutional output is passed through an adaptive average pooling layer, reducing the spatial resolution to  $1 \times 1$  and producing a compact 64-dimensional feature vector. This eliminates the need to flatten large tensors, significantly reducing memory requirements.

The classifier includes a fully connected layer mapping 64 to 128 dimensions, followed by ReLU and dropout ( $p=0.3$ ), and a final layer mapping to 3 output classes. The model is trained using the AdamW optimizer and cross-entropy loss.

---

## (b) Training & Validation Code

```
In [29]: !unzip -q "garment_images.zip" -d garment_images  
replace garment_images/train_labels.csv? [y]es, [n]o, [A]ll, [N]one, [r]ename: A
```

```
In [30]: # Load the files  
import torch, os  
import torch.nn as nn  
import torch.nn.functional as F  
import numpy as np  
from torchvision import datasets, transforms  
from torch.utils.data import DataLoader, random_split  
from torch.optim import AdamW  
from sklearn.metrics import accuracy_score  
  
# Data Transformation  
train_tf = transforms.Compose([  
    transforms.RandomHorizontalFlip(),  
    transforms.RandomRotation(15),  
    transforms.ColorJitter(0.2, 0.2),  
    transforms.ToTensor(),  
    transforms.Normalize([0.5]*3, [0.5]*3)  
)  
val_tf = transforms.Compose([
```

```

        transforms.ToTensor(),
        transforms.Normalize([0.5]*3, [0.5]*3)
    ])

# Data Loading
data_path = './garment_images' # Place folder in Colab root
full_ds = datasets.ImageFolder(data_path, transform=train_tf)
train_ds, val_ds = random_split(full_ds, [int(0.8*len(full_ds)), len(full_ds)-int(0.8*len(full_ds))])
val_ds.dataset.transform = val_tf

train_dl = DataLoader(train_ds, batch_size=32, shuffle=True)
val_dl = DataLoader(val_ds, batch_size=32)

# Defining the model
class GarmentNet(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(3, 16, 3, padding=1), nn.BatchNorm2d(16), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(16, 32, 3, padding=1), nn.BatchNorm2d(32), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(32, 64, 3, padding=1), nn.BatchNorm2d(64), nn.ReLU(), nn.MaxPool2d(2),
            nn.Conv2d(64, 64, 3, padding=1), nn.BatchNorm2d(64), nn.ReLU(), nn.AdaptiveAvgPool2d(1)
        )
        self.fc = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64, 128), nn.ReLU(), nn.Dropout(0.3),
            nn.Linear(128, 3)
        )

    def forward(self, x):
        return self.fc(self.conv(x))

# Training
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = GarmentNet().to(device)
opt = AdamW(model.parameters(), lr=1e-3, weight_decay=1e-4)
loss_fn = nn.CrossEntropyLoss()

os.makedirs("garment_model", exist_ok=True)
best_acc = 0.0

for epoch in range(1, 21):
    model.train()
    for xb, yb in train_dl:
        xb, yb = xb.to(device), yb.to(device)
        opt.zero_grad()
        loss = loss_fn(model(xb), yb)
        loss.backward()
        opt.step()

    model.eval()
    preds, trues = [], []
    with torch.no_grad():
        for xb, yb in val_dl:
            xb = xb.to(device)
            out = model(xb).cpu().argmax(dim=1).numpy()

```

```

        preds.append(out)
        trues.append(yb.numpy())
preds = np.concatenate(preds)
trues = np.concatenate(trues)
acc = accuracy_score(trues, preds) * 100
print(f"Epoch {epoch:02d} - Val Accuracy: {acc:.2f}%")
if acc > best_acc:
    best_acc = acc
    torch.save(model.state_dict(), "garment_model/net_weights.pth")

print(f"\n\n{checkmark} Best Val Accuracy: {best_acc:.2f}%")
print(f"{checkmark} Weight file size: {os.path.getsize("garment_model/net_weights.pth"):.2f} MiB")

```

Epoch 01 – Val Accuracy: 69.77%  
 Epoch 02 – Val Accuracy: 56.27%  
 Epoch 03 – Val Accuracy: 50.19%  
 Epoch 04 – Val Accuracy: 62.93%  
 Epoch 05 – Val Accuracy: 74.71%  
 Epoch 06 – Val Accuracy: 89.16%  
 Epoch 07 – Val Accuracy: 82.51%  
 Epoch 08 – Val Accuracy: 84.98%  
 Epoch 09 – Val Accuracy: 63.50%  
 Epoch 10 – Val Accuracy: 58.37%  
 Epoch 11 – Val Accuracy: 60.08%  
 Epoch 12 – Val Accuracy: 82.32%  
 Epoch 13 – Val Accuracy: 83.65%  
 Epoch 14 – Val Accuracy: 93.35%  
 Epoch 15 – Val Accuracy: 87.83%  
 Epoch 16 – Val Accuracy: 90.87%  
 Epoch 17 – Val Accuracy: 83.84%  
 Epoch 18 – Val Accuracy: 89.73%  
 Epoch 19 – Val Accuracy: 79.28%  
 Epoch 20 – Val Accuracy: 86.69%

Best Val Accuracy: 93.35%  
 Weight file size: 0.28 MiB

Write predict\_class.py

This inference module defines a compact convolutional neural network for classifying 256×256 RGB garment images into three categories.

It includes:

1. Standard normalization
2. Loading pretrained weights
3. Returns class predictions

```

In [6]: import torch
import torch.nn as nn
import torch.nn.functional as F

class GarmentNet(nn.Module):
    def __init__(self):
        super().__init__()

```

```

        self.conv = nn.Sequential(
            nn.Conv2d(3, 16, 3, padding=1), nn.BatchNorm2d(16), nn.ReLU(), nn.MaxPool2d(2, 2),
            nn.Conv2d(16, 32, 3, padding=1), nn.BatchNorm2d(32), nn.ReLU(), nn.MaxPool2d(2, 2),
            nn.Conv2d(32, 64, 3, padding=1), nn.BatchNorm2d(64), nn.ReLU(), nn.MaxPool2d(2, 2),
            nn.Conv2d(64, 64, 3, padding=1), nn.BatchNorm2d(64), nn.ReLU(), nn.AdaptiveAvgPool2d(1)
        )
        self.fc = nn.Sequential(
            nn.Flatten(),
            nn.Linear(64, 128), nn.ReLU(), nn.Dropout(0.3),
            nn.Linear(128, 3)
        )

    def forward(self, x):
        return self.fc(self.conv(x))

    _net = GarmentNet()
    _net.load_state_dict(torch.load("garment_model/net_weights.pth", map_location="cpu"))
    _net.eval()

    def predict(images):
        """ images: B x 3 x 256 x 256 with values in (0, 1) """
        mean = torch.tensor([0.5, 0.5, 0.5]).view(1, 3, 1, 1)
        std = torch.tensor([0.5, 0.5, 0.5]).view(1, 3, 1, 1)
        x = (images - mean) / std
        with torch.no_grad():
            return torch.argmax(_net(x), dim=1, keepdim=True)

```

In [7]:

```

import os

weight_path = "garment_model/net_weights.pth"
size_mb = os.path.getsize(weight_path) / (1024 * 1024)
print(f"Model weight size: {size_mb:.2f} MiB")

```

Model weight size: 0.28 MiB

## Question 6: Neural Image Compression via Autoencoder

(a) A convolutional autoencoder was designed to compress  $192 \times 160$  grayscale facial images into a compact 16-dimensional latent representation and reconstruct them with high perceptual quality.

The encoder consists of four strided convolutional layers, which progressively reduce spatial resolution from  $192 \times 160$  to  $12 \times 10$ , while increasing the feature depth. This setup effectively captures hierarchical spatial features. A fully connected bottleneck layer then reduces the resulting feature map ( $256 \times 12 \times 10$ ) into a 16D latent vector, achieving strong compression.

The decoder mirrors this structure in reverse. It uses a fully connected layer to reshape the 16D latent code into a  $256 \times 12 \times 10$  tensor, followed by four ConvTranspose2d layers that upsample it back to the original resolution of  $192 \times 160$ . ReLU activations are used

throughout the layers to introduce non-linearity, and a final sigmoid activation ensures output pixel values are within the range [0, 1].

The architecture avoids batch normalization and skip connections, keeping the model lightweight and interpretable while achieving high reconstruction fidelity.

(b) The model was trained to minimize mean squared error (MSE) loss between the input and reconstructed image. In addition, structural similarity index (SSIM) was monitored on a 20% validation split to assess perceptual quality.

The AdamW optimizer was used with a learning rate of 5e-4 and weight decay of 1e-5 to promote stable and generalizable learning. A ReduceLROnPlateau scheduler was employed to dynamically reduce the learning rate when validation SSIM plateaued, encouraging continued improvement without overfitting.

The training and validation MSE loss curves show steady decrease, while SSIM values rise consistently, reaching SSIM  $\approx 0.7853$  by epoch 100. This indicates that the network is learning to reconstruct visually accurate images and that the chosen hyperparameters (optimizer, scheduler, learning rate, weight decay) effectively support convergence and generalization.

```
In [13]: # Cell 1: Setup & Data
!unzip -q "face_images.zip" -d faces_data

replace faces_data/0001.jpg? [y]es, [n]o, [A]ll, [N]one, [r]ename: A
```

```
In [26]: # train_compression.py

import os, glob
import numpy as np
import matplotlib.pyplot as plt
from PIL import Image

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader, random_split
from torchvision import transforms
from skimage.metrics import structural_similarity as ssim

# Dataset
class FacesDataset(Dataset):
    def __init__(self, files, tf):
        self.files = files
        self.tf = tf
    def __len__(self): return len(self.files)
    def __getitem__(self, i):
        img = Image.open(self.files[i]).convert('L')
        return self.tf(img)

# Encoder
class Encoder(nn.Module):
```

```

def __init__(self, latent_dim=16):
    super().__init__()
    self.conv = nn.Sequential(
        nn.Conv2d(1, 32, 4, 2, 1), nn.ReLU(True),
        nn.Conv2d(32, 64, 4, 2, 1), nn.ReLU(True),
        nn.Conv2d(64, 128, 4, 2, 1), nn.ReLU(True),
        nn.Conv2d(128, 256, 4, 2, 1), nn.ReLU(True),
    )
    self.fc = nn.Linear(256*12*10, latent_dim)
def forward(self, x):
    x = self.conv(x)
    x = x.view(x.size(0), -1)
    return self.fc(x)

# Decoder
class Decoder(nn.Module):
    def __init__(self, latent_dim=16):
        super().__init__()
        self.fc = nn.Linear(latent_dim, 256*12*10)
        self.deconv = nn.Sequential(
            nn.ConvTranspose2d(256, 128, 4, 2, 1), nn.ReLU(True),
            nn.ConvTranspose2d(128, 64, 4, 2, 1), nn.ReLU(True),
            nn.ConvTranspose2d(64, 32, 4, 2, 1), nn.ReLU(True),
            nn.ConvTranspose2d(32, 1, 4, 2, 1), nn.Sigmoid()
        )
    def forward(self, z):
        x = self.fc(z).view(z.size(0), 256, 12, 10)
        return self.deconv(x)

# Setup
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
enc = Encoder().to(device)
dec = Decoder().to(device)
params = list(enc.parameters()) + list(dec.parameters())

opt = optim.AdamW(params, lr=5e-4, weight_decay=1e-5)
sched = optim.lr_scheduler.ReduceLROnPlateau(opt, mode='max', factor=0.5, patience=10)
mse_loss = nn.MSELoss()

# Dataset and DataLoader
files = sorted(glob.glob('faces_data/*.jpg'))
tf = transforms.ToTensor()
ds_all = FacesDataset(files, tf)
n_val = int(0.2 * len(ds_all))
train_ds, val_ds = random_split(ds_all, [len(ds_all)-n_val, n_val])
train_dl = DataLoader(train_ds, batch_size=32, shuffle=True)
val_dl = DataLoader(val_ds, batch_size=32, shuffle=False)

# Training Loop
train_losses, val_losses, val_ssims = [], [], []
best_ssim, best_ep = -1, 0

for epoch in range(1, 101):
    enc.train(); dec.train()
    tr_loss = 0.0
    for xb in train_dl:

```

```

        xb = xb.to(device)
        z = enc(xb)
        xr = dec(z)
        loss = mse_loss(xr, xb)
        opt.zero_grad(); loss.backward(); opt.step()
        tr_loss += loss.item() * xb.size(0)
        train_losses.append(tr_loss / len(train_ds))

        enc.eval(); dec.eval()
        va_loss, ssim_sum, count = 0.0, 0.0, 0
        with torch.no_grad():
            for xb in val_dl:
                xb = xb.to(device)
                z = enc(xb)
                xr = dec(z)
                va_loss += mse_loss(xr, xb).item() * xb.size(0)
                xb_np, xr_np = xb.cpu().numpy(), xr.cpu().numpy()
                for i in range(xb_np.shape[0]):
                    ssim_sum += ssim(xb_np[i,0], xr_np[i,0], data_range=1.0)
                    count += 1
            val_losses.append(va_loss / len(val_ds))
            avg_ssim = ssim_sum / count
            val_ssims.append(avg_ssim)
            sched.step(avg_ssim)

        if avg_ssim > best_ssim:
            best_ssim, best_ep = avg_ssim, epoch
            os.makedirs('compression_model', exist_ok=True)
            torch.save(enc.state_dict(), 'compression_model/encoder.pth')
            torch.save(dec.state_dict(), 'compression_model/decoder.pth')

        print(f"Epoch {epoch:03d} | Tr MSE: {train_losses[-1]:.6f} | "
              f"Va MSE: {val_losses[-1]:.6f} | Va SSIM: {avg_ssim:.4f}")

    print(f"\nBest Val SSIM = {best_ssim:.4f} @ epoch {best_ep}")

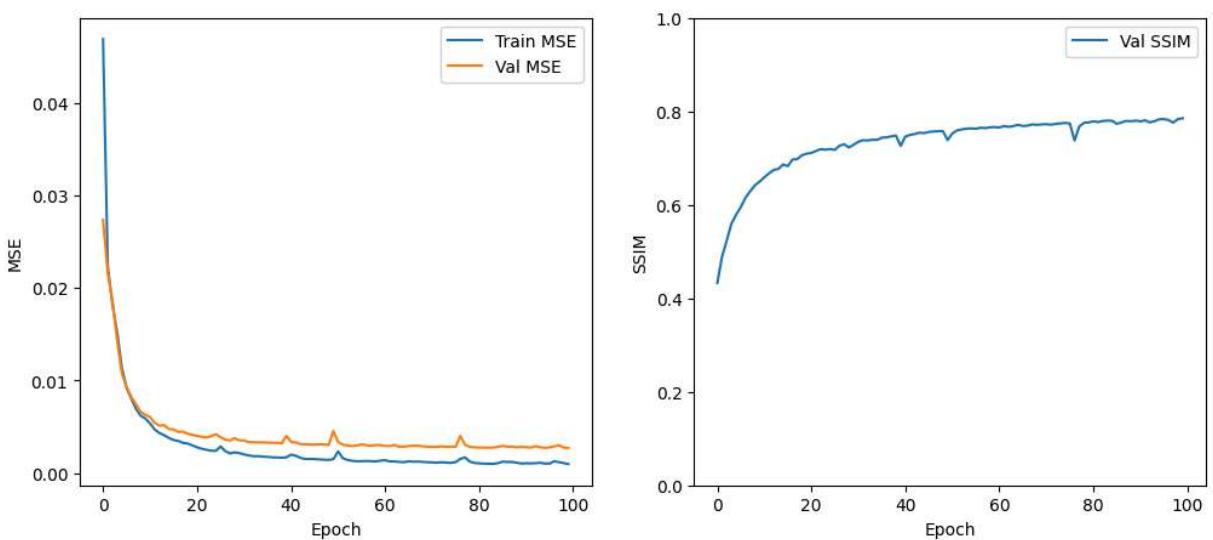
# Plot
plt.figure(figsize=(12,5))
plt.subplot(1,2,1)
plt.plot(train_losses, label='Train MSE')
plt.plot(val_losses, label='Val MSE')
plt.xlabel('Epoch'); plt.ylabel('MSE'); plt.legend()
plt.subplot(1,2,2)
plt.plot(val_ssims, label='Val SSIM')
plt.xlabel('Epoch'); plt.ylabel('SSIM'); plt.ylim(0,1.0); plt.legend()
plt.show()

```

Epoch 001	Tr MSE: 0.046852	Va MSE: 0.027341	Va SSIM: 0.4329
Epoch 002	Tr MSE: 0.022287	Va MSE: 0.021623	Va SSIM: 0.4887
Epoch 003	Tr MSE: 0.018209	Va MSE: 0.018592	Va SSIM: 0.5235
Epoch 004	Tr MSE: 0.015189	Va MSE: 0.014309	Va SSIM: 0.5597
Epoch 005	Tr MSE: 0.011355	Va MSE: 0.010754	Va SSIM: 0.5787
Epoch 006	Tr MSE: 0.009227	Va MSE: 0.009299	Va SSIM: 0.5953
Epoch 007	Tr MSE: 0.008050	Va MSE: 0.008210	Va SSIM: 0.6152
Epoch 008	Tr MSE: 0.006956	Va MSE: 0.007430	Va SSIM: 0.6292
Epoch 009	Tr MSE: 0.006195	Va MSE: 0.006615	Va SSIM: 0.6418
Epoch 010	Tr MSE: 0.005895	Va MSE: 0.006295	Va SSIM: 0.6496
Epoch 011	Tr MSE: 0.005364	Va MSE: 0.006044	Va SSIM: 0.6589
Epoch 012	Tr MSE: 0.004703	Va MSE: 0.005411	Va SSIM: 0.6673
Epoch 013	Tr MSE: 0.004332	Va MSE: 0.005134	Va SSIM: 0.6747
Epoch 014	Tr MSE: 0.004090	Va MSE: 0.005179	Va SSIM: 0.6769
Epoch 015	Tr MSE: 0.003801	Va MSE: 0.004765	Va SSIM: 0.6864
Epoch 016	Tr MSE: 0.003575	Va MSE: 0.004722	Va SSIM: 0.6828
Epoch 017	Tr MSE: 0.003479	Va MSE: 0.004446	Va SSIM: 0.6971
Epoch 018	Tr MSE: 0.003262	Va MSE: 0.004471	Va SSIM: 0.6978
Epoch 019	Tr MSE: 0.003207	Va MSE: 0.004257	Va SSIM: 0.7060
Epoch 020	Tr MSE: 0.002995	Va MSE: 0.004133	Va SSIM: 0.7096
Epoch 021	Tr MSE: 0.002788	Va MSE: 0.004028	Va SSIM: 0.7110
Epoch 022	Tr MSE: 0.002635	Va MSE: 0.003916	Va SSIM: 0.7153
Epoch 023	Tr MSE: 0.002517	Va MSE: 0.003832	Va SSIM: 0.7190
Epoch 024	Tr MSE: 0.002416	Va MSE: 0.003997	Va SSIM: 0.7180
Epoch 025	Tr MSE: 0.002409	Va MSE: 0.004201	Va SSIM: 0.7194
Epoch 026	Tr MSE: 0.002878	Va MSE: 0.003881	Va SSIM: 0.7176
Epoch 027	Tr MSE: 0.002374	Va MSE: 0.003610	Va SSIM: 0.7265
Epoch 028	Tr MSE: 0.002129	Va MSE: 0.003530	Va SSIM: 0.7296
Epoch 029	Tr MSE: 0.002243	Va MSE: 0.003773	Va SSIM: 0.7226
Epoch 030	Tr MSE: 0.002165	Va MSE: 0.003532	Va SSIM: 0.7285
Epoch 031	Tr MSE: 0.002018	Va MSE: 0.003521	Va SSIM: 0.7348
Epoch 032	Tr MSE: 0.001913	Va MSE: 0.003342	Va SSIM: 0.7384
Epoch 033	Tr MSE: 0.001824	Va MSE: 0.003336	Va SSIM: 0.7377
Epoch 034	Tr MSE: 0.001831	Va MSE: 0.003317	Va SSIM: 0.7393
Epoch 035	Tr MSE: 0.001776	Va MSE: 0.003319	Va SSIM: 0.7388
Epoch 036	Tr MSE: 0.001753	Va MSE: 0.003290	Va SSIM: 0.7438
Epoch 037	Tr MSE: 0.001705	Va MSE: 0.003274	Va SSIM: 0.7442
Epoch 038	Tr MSE: 0.001695	Va MSE: 0.003275	Va SSIM: 0.7465
Epoch 039	Tr MSE: 0.001659	Va MSE: 0.003205	Va SSIM: 0.7483
Epoch 040	Tr MSE: 0.001703	Va MSE: 0.004029	Va SSIM: 0.7261
Epoch 041	Tr MSE: 0.001994	Va MSE: 0.003358	Va SSIM: 0.7460
Epoch 042	Tr MSE: 0.001884	Va MSE: 0.003309	Va SSIM: 0.7496
Epoch 043	Tr MSE: 0.001637	Va MSE: 0.003130	Va SSIM: 0.7515
Epoch 044	Tr MSE: 0.001529	Va MSE: 0.003138	Va SSIM: 0.7544
Epoch 045	Tr MSE: 0.001524	Va MSE: 0.003088	Va SSIM: 0.7534
Epoch 046	Tr MSE: 0.001521	Va MSE: 0.003072	Va SSIM: 0.7560
Epoch 047	Tr MSE: 0.001478	Va MSE: 0.003105	Va SSIM: 0.7569
Epoch 048	Tr MSE: 0.001449	Va MSE: 0.003096	Va SSIM: 0.7574
Epoch 049	Tr MSE: 0.001429	Va MSE: 0.003039	Va SSIM: 0.7577
Epoch 050	Tr MSE: 0.001502	Va MSE: 0.004558	Va SSIM: 0.7385
Epoch 051	Tr MSE: 0.002338	Va MSE: 0.003369	Va SSIM: 0.7527
Epoch 052	Tr MSE: 0.001621	Va MSE: 0.003075	Va SSIM: 0.7593
Epoch 053	Tr MSE: 0.001417	Va MSE: 0.002985	Va SSIM: 0.7613
Epoch 054	Tr MSE: 0.001331	Va MSE: 0.002922	Va SSIM: 0.7628
Epoch 055	Tr MSE: 0.001288	Va MSE: 0.002986	Va SSIM: 0.7637
Epoch 056	Tr MSE: 0.001287	Va MSE: 0.003103	Va SSIM: 0.7628

Epoch 057	Tr MSE: 0.001311	Va MSE: 0.003013	Va SSIM: 0.7650
Epoch 058	Tr MSE: 0.001291	Va MSE: 0.002950	Va SSIM: 0.7641
Epoch 059	Tr MSE: 0.001271	Va MSE: 0.003032	Va SSIM: 0.7657
Epoch 060	Tr MSE: 0.001346	Va MSE: 0.003010	Va SSIM: 0.7664
Epoch 061	Tr MSE: 0.001391	Va MSE: 0.002939	Va SSIM: 0.7653
Epoch 062	Tr MSE: 0.001268	Va MSE: 0.002915	Va SSIM: 0.7687
Epoch 063	Tr MSE: 0.001261	Va MSE: 0.003051	Va SSIM: 0.7667
Epoch 064	Tr MSE: 0.001212	Va MSE: 0.002850	Va SSIM: 0.7687
Epoch 065	Tr MSE: 0.001179	Va MSE: 0.002861	Va SSIM: 0.7712
Epoch 066	Tr MSE: 0.001274	Va MSE: 0.002918	Va SSIM: 0.7687
Epoch 067	Tr MSE: 0.001233	Va MSE: 0.002961	Va SSIM: 0.7694
Epoch 068	Tr MSE: 0.001245	Va MSE: 0.002949	Va SSIM: 0.7721
Epoch 069	Tr MSE: 0.001212	Va MSE: 0.002897	Va SSIM: 0.7712
Epoch 070	Tr MSE: 0.001179	Va MSE: 0.002848	Va SSIM: 0.7720
Epoch 071	Tr MSE: 0.001163	Va MSE: 0.002820	Va SSIM: 0.7726
Epoch 072	Tr MSE: 0.001127	Va MSE: 0.002831	Va SSIM: 0.7714
Epoch 073	Tr MSE: 0.001181	Va MSE: 0.002891	Va SSIM: 0.7734
Epoch 074	Tr MSE: 0.001133	Va MSE: 0.002832	Va SSIM: 0.7742
Epoch 075	Tr MSE: 0.001116	Va MSE: 0.002861	Va SSIM: 0.7753
Epoch 076	Tr MSE: 0.001188	Va MSE: 0.002839	Va SSIM: 0.7740
Epoch 077	Tr MSE: 0.001533	Va MSE: 0.004012	Va SSIM: 0.7377
Epoch 078	Tr MSE: 0.001696	Va MSE: 0.003067	Va SSIM: 0.7675
Epoch 079	Tr MSE: 0.001254	Va MSE: 0.002839	Va SSIM: 0.7755
Epoch 080	Tr MSE: 0.001096	Va MSE: 0.002783	Va SSIM: 0.7764
Epoch 081	Tr MSE: 0.001061	Va MSE: 0.002759	Va SSIM: 0.7786
Epoch 082	Tr MSE: 0.001018	Va MSE: 0.002756	Va SSIM: 0.7771
Epoch 083	Tr MSE: 0.001005	Va MSE: 0.002741	Va SSIM: 0.7794
Epoch 084	Tr MSE: 0.001001	Va MSE: 0.002761	Va SSIM: 0.7801
Epoch 085	Tr MSE: 0.001069	Va MSE: 0.002839	Va SSIM: 0.7798
Epoch 086	Tr MSE: 0.001240	Va MSE: 0.002937	Va SSIM: 0.7733
Epoch 087	Tr MSE: 0.001204	Va MSE: 0.002857	Va SSIM: 0.7763
Epoch 088	Tr MSE: 0.001212	Va MSE: 0.002869	Va SSIM: 0.7796
Epoch 089	Tr MSE: 0.001127	Va MSE: 0.002787	Va SSIM: 0.7788
Epoch 090	Tr MSE: 0.001033	Va MSE: 0.002836	Va SSIM: 0.7802
Epoch 091	Tr MSE: 0.001072	Va MSE: 0.002802	Va SSIM: 0.7785
Epoch 092	Tr MSE: 0.001051	Va MSE: 0.002761	Va SSIM: 0.7809
Epoch 093	Tr MSE: 0.001073	Va MSE: 0.002911	Va SSIM: 0.7764
Epoch 094	Tr MSE: 0.001121	Va MSE: 0.002790	Va SSIM: 0.7791
Epoch 095	Tr MSE: 0.001030	Va MSE: 0.002724	Va SSIM: 0.7829
Epoch 096	Tr MSE: 0.001033	Va MSE: 0.002781	Va SSIM: 0.7835
Epoch 097	Tr MSE: 0.001283	Va MSE: 0.002886	Va SSIM: 0.7815
Epoch 098	Tr MSE: 0.001185	Va MSE: 0.003014	Va SSIM: 0.7760
Epoch 099	Tr MSE: 0.001089	Va MSE: 0.002761	Va SSIM: 0.7834
Epoch 100	Tr MSE: 0.000979	Va MSE: 0.002696	Va SSIM: 0.7853

Best Val SSIM = 0.7853 @ epoch 100



```
In [24]: import os; os.makedirs('compression_model', exist_ok=True)
torch.save(enc.state_dict(), 'compression_model/encoder.pth')
torch.save(dec.state_dict(), 'compression_model/decoder.pth')
```

compress\_images.py

```
In [25]: # compress_images.py

import torch
import torch.nn as nn

class Encoder(nn.Module):
    def __init__(self, latent_dim=16):
        super().__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(1, 32, 4, 2, 1), nn.ReLU(True),
            nn.Conv2d(32, 64, 4, 2, 1), nn.ReLU(True),
            nn.Conv2d(64, 128, 4, 2, 1), nn.ReLU(True),
            nn.Conv2d(128, 256, 4, 2, 1), nn.ReLU(True),
        )
        self.fc = nn.Linear(256*12*10, latent_dim)
    def forward(self, x):
        x = self.conv(x)
        x = x.view(x.size(0), -1)
        return self.fc(x)

class Decoder(nn.Module):
    def __init__(self, latent_dim=16):
        super().__init__()
        self.fc = nn.Linear(latent_dim, 256*12*10)
        self.deconv = nn.Sequential(
            nn.ConvTranspose2d(256, 128, 4, 2, 1), nn.ReLU(True),
            nn.ConvTranspose2d(128, 64, 4, 2, 1), nn.ReLU(True),
            nn.ConvTranspose2d(64, 32, 4, 2, 1), nn.ReLU(True),
            nn.ConvTranspose2d(32, 1, 4, 2, 1), nn.Sigmoid()
        )
    def forward(self, z):
        x = self.fc(z).view(z.size(0), 256, 12, 10)
```

```

        return self.deconv(x)

# Load model
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
_enc = Encoder().to(device)
_dec = Decoder().to(device)

_enc.load_state_dict(torch.load('compression_model/encoder.pth', map_location=device))
_dec.load_state_dict(torch.load('compression_model/decoder.pth', map_location=device))
_enc.eval()
_dec.eval()

def encode(images: torch.Tensor) -> torch.Tensor:
    with torch.no_grad():
        return _enc(images.to(device)).cpu()

def decode(latents: torch.Tensor) -> torch.Tensor:
    with torch.no_grad():
        return _dec(latents.to(device)).cpu()

```

## ZIP all models and inference files

In [31]: `!zip -r models_inference.zip compression_model waist_model garment_model`

```

adding: compression_model/ (stored 0%)
adding: compression_model/encoder.pth (deflated 7%)
adding: compression_model/decoder.pth (deflated 7%)
adding: waist_model/ (stored 0%)
adding: waist_model/net_weights_fold3.pth (deflated 31%)
adding: waist_model/net_weights_fold1.pth (deflated 31%)
adding: waist_model/net_weights_fold4.pth (deflated 31%)
adding: waist_model/net_weights_fold5.pth (deflated 31%)
adding: waist_model/scaler_y.joblib (deflated 31%)
adding: waist_model/scaler_X.joblib (deflated 21%)
adding: waist_model/net_weights_fold2.pth (deflated 31%)
adding: garment_model/ (stored 0%)
adding: garment_model/net_weights.pth (deflated 10%)

```

In [32]: `from google.colab import files  
files.download('models_inference.zip')`

In [33]: `from google.colab import files  
files.download('PADL-Q3-result.txt')`